

# 5. Transformers

## *Computational Music Creativity*



Universitat  
Pompeu Fabra  
*Barcelona*

**MTG**  
Music Technology  
Group



**NEW,  
SHINY AI MODEL**

**ANY AI  
MUSIC ENGINEER**


**60+ YEARS  
OF SOLID  
TECHNIQUES**







# Real-time scores



## A reference problem

---

My friends like tomatoes because they are tasty



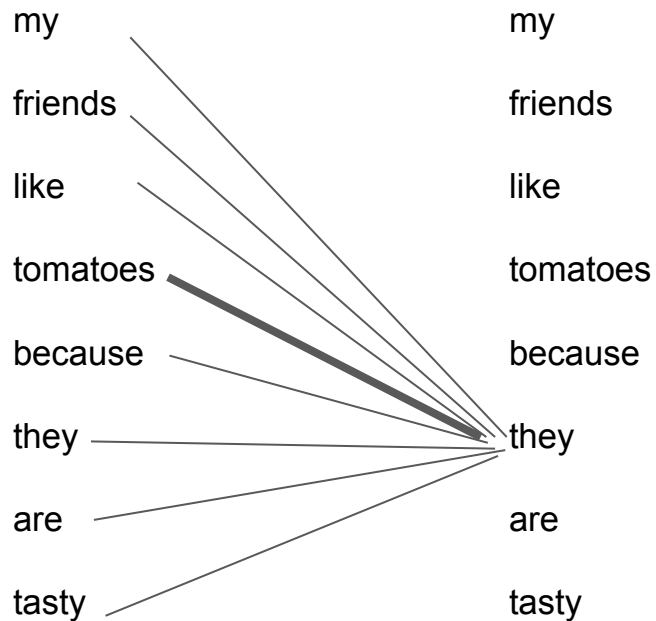
## A reference problem

---

My friends like tomatoes because **they** are tasty

# Self-attention: Intuition

---



# What matrices do we have in self-attention?

---

# Query, key, value matrices

---

Query (Q)

I	1.3	0.8
like	0.7	3.5
cats	1.9	0.1

Key (K)

I	0.6	2.4
like	0.8	1.7
cats	2.5	0.3

Value (V)

I	0.4	1.0
like	1.2	2.8
cats	1.7	0.2

How do we derive Q, K, V?

---

# How do we derive Q, K, V?

---

- Multiply input matrix by 3 weight matrices
- Learn weights during training

$$IW_Q = Q$$

$$IW_K = K$$

$$IW_V = V$$

# Self-attention: Formalisation

---

$$Z(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

## Self-attention: Step 1

---

$$Z(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$



# Self-attention: Step 1

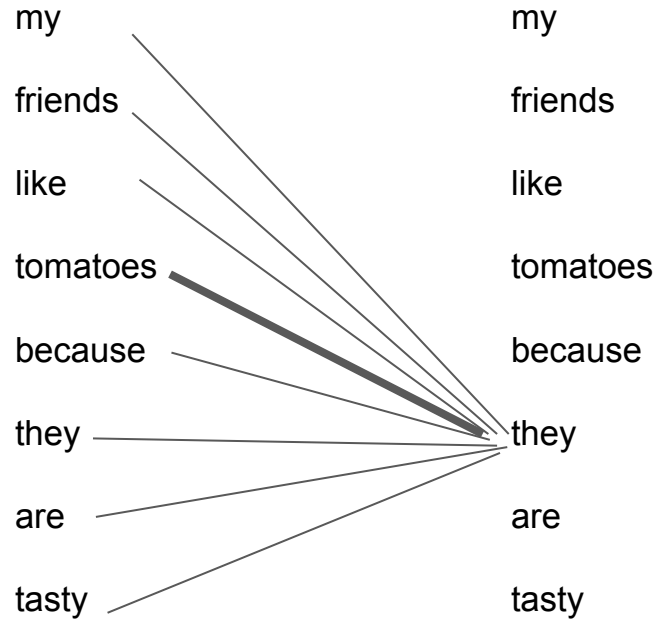
---

$$QK^T = \begin{matrix} & \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} & \begin{bmatrix} 1.3 & 0.8 \\ 0.7 & 3.5 \\ 1.9 & 0.1 \end{bmatrix} & \begin{matrix} q_1 \\ q_2 \\ q_3 \end{matrix} \end{matrix} \begin{matrix} \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} & \begin{matrix} k_1 \\ k_2 \\ k_3 \end{matrix} \\ \begin{matrix} \text{I} & \text{like} & \text{cats} \end{matrix} & \begin{bmatrix} 0.6 & 0.8 & 2.5 \\ 2.4 & 1.7 & 0.3 \end{bmatrix} \end{matrix} = \begin{bmatrix} q_1 k_1 & q_1 k_2 & q_1 k_3 \\ q_2 k_1 & q_2 k_2 & q_2 k_3 \\ q_3 k_1 & q_3 k_2 & q_3 k_3 \end{bmatrix} = \begin{matrix} & \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} & \begin{bmatrix} 2.7 & 2.4 & 3.49 \\ 8.82 & 6.51 & 2.8 \\ 1.38 & 1.69 & 4.78 \end{bmatrix} \end{matrix}$$

$Q$   $K^T$

# What are Q and K really?

---




## Self-attention: Step 2

---

$$Z(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

# What are Q and K really?

---



## Self-attention: Step 3

---

$$Z(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

# Self-attention: Step 3

---

- Normalize similarity scores
- Apply *softmax*
- Each word vector (row) adds up to 1 (probability)

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) = \begin{matrix} & \text{I} & \text{like} & \text{cats} \\ \text{I} & 0.7 & 0.2 & 0.1 \\ \text{like} & 0.2 & 0.6 & 0.2 \\ \text{cats} & 0.4 & 0.1 & 0.5 \end{matrix}$$

\*values in the matrix completely made up

## Self-attention: Step 3

---

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right)$$

## Self-attention: Step 3

---

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right)$$

Attention score

Relevance of different parts of the sequence to each other



## Self-attention: Step 4

---

$$Z(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

## Self-attention: Step 4

---

$$Z = \begin{matrix} & \text{I} & \text{like} & \text{cats} \\ \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} & \begin{bmatrix} 0.7 & 0.2 & 0.1 \\ 0.2 & 0.6 & 0.2 \\ 0.4 & 0.1 & 0.5 \end{bmatrix} \end{matrix} \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} \begin{bmatrix} 0.4 & 1.0 \\ 1.2 & 2.8 \\ 1.7 & 0.2 \end{bmatrix} \begin{matrix} \mathbf{v}_1 \\ \mathbf{v}_2 \\ \mathbf{v}_3 \end{matrix} = \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} \begin{bmatrix} 0.69 & 1.28 \\ 1.14 & 1.92 \\ 1.13 & 0.78 \end{bmatrix} = \begin{bmatrix} \vec{z}_1 \\ \vec{z}_2 \\ \vec{z}_3 \end{bmatrix}$$

$\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)$        $V$

# Self-attention for word “I”

---

$$Z = \begin{matrix} & \text{I} & \text{like} & \text{cats} \\ \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} & \begin{bmatrix} 0.7 & 0.2 & 0.1 \\ 0.2 & 0.6 & 0.2 \\ 0.4 & 0.1 & 0.5 \end{bmatrix} \end{matrix} \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} \begin{bmatrix} 0.4 & 1.0 \\ 1.2 & 2.8 \\ 1.7 & 0.2 \end{bmatrix} \begin{matrix} \vec{v}_1 \\ \vec{v}_2 \\ \vec{v}_3 \end{matrix} = \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} \begin{bmatrix} 0.69 & 1.28 \\ 1.14 & 1.92 \\ 1.13 & 0.78 \end{bmatrix} = \begin{bmatrix} \vec{z}_1 \\ \vec{z}_2 \\ \vec{z}_3 \end{bmatrix}$$

$\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) \quad V$

$$\vec{z}_1 = 0.7\vec{v}_1 + 0.2\vec{v}_2 + 0.1\vec{v}_3 = 0.7 \begin{bmatrix} 0.4 & 1.0 \end{bmatrix} + 0.2 \begin{bmatrix} 1.2 & 2.8 \end{bmatrix} + 0.1 \begin{bmatrix} 1.7 & 0.2 \end{bmatrix}$$

I                      like                      cats                      I                      like                      cats

Sum of the value vectors weighted by the scores

## A reference problem: Solved

---

My friends like tomatoes because they are tasty

$$\vec{z}_{they} = 0.0\vec{v}_1 + 0.0\vec{v}_2 + 0.0\vec{v}_3 + 0.9\vec{v}_4 + 0.0\vec{v}_5 + 0.1\vec{v}_6 + 0.0\vec{v}_7 + 0.0\vec{v}_8$$

my                  friends                  like                  tomatoes                  because                  they                  are                  tasty

# What's multi-head attention?

---

# What's multi-head attention?

---

- Run multiple instances of the self-attention mechanism in parallel
- Compute as many Q, K, V, Z matrices as the number of heads

$$Z = \text{concatenate}(Z_1, Z_2, Z_3, \dots, Z_n)W_0$$

**WHY MULTIPLE HEADS?**



# Why positional encoding?

---



## Positional encoding: Strategy

---

$$I' = \underbrace{\begin{bmatrix} 0.2 & 1.2 \\ 0.5 & 4.1 \\ 2.1 & 0.4 \end{bmatrix}}_I + \underbrace{\begin{bmatrix} 0.5 & 1.0 \\ 2.5 & 1.3 \\ 1.1 & 0.3 \end{bmatrix}}_P = \begin{bmatrix} 0.7 & 2.2 \\ 3.0 & 5.4 \\ 3.2 & 0.7 \end{bmatrix}$$

# How do we compute $P$ ?

---

# How do we compute P?

---

$$P(pos, 2i) = \sin \left( \frac{pos}{10000^{2i/dimension_{model}}} \right)$$

$$P(pos, 2i + 1) = \cos \left( \frac{pos}{10000^{2i/dimension_{model}}} \right)$$

# How do we compute P?

---

$$P(pos, 2i) = \sin\left(\frac{pos}{10000^{2i/dimension_{model}}}\right)$$


$$P(pos, 2i + 1) = \cos\left(\frac{pos}{10000^{2i/dimension_{model}}}\right)$$

$$P = \begin{matrix} \text{Spaghetti} \\ \text{monster} \\ \text{is} \\ \text{great} \end{matrix} \begin{bmatrix} \sin\left(\frac{0}{10000^{2.0/3}}\right) & \cos\left(\frac{0}{10000^{2.1/2}}\right) & \sin\left(\frac{0}{10000^{2.2/3}}\right) \\ \sin\left(\frac{1}{10000^{2.0/3}}\right) & \cos\left(\frac{1}{10000^{2.1/2}}\right) & \sin\left(\frac{1}{10000^{2.2/3}}\right) \\ \sin\left(\frac{2}{10000^{2.0/3}}\right) & \cos\left(\frac{2}{10000^{2.1/2}}\right) & \sin\left(\frac{2}{10000^{2.2/3}}\right) \\ \sin\left(\frac{3}{10000^{2.0/3}}\right) & \cos\left(\frac{3}{10000^{2.1/2}}\right) & \sin\left(\frac{3}{10000^{2.2/3}}\right) \end{bmatrix}$$

# Other components missing from encoder?


---


- Feed-forward
- Add & Norm




# Other components missing from encoder?

---



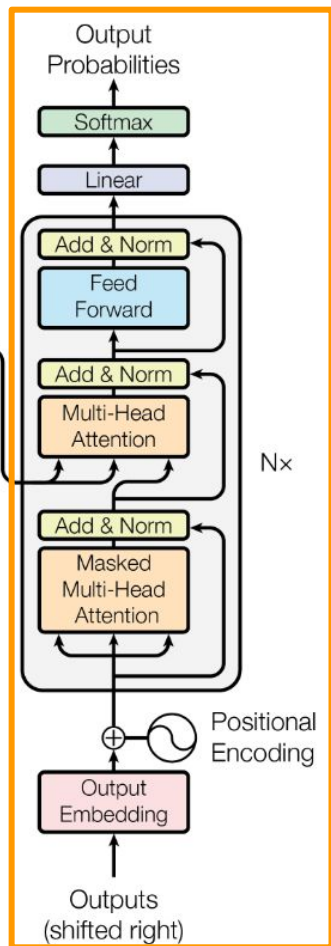


## Encoder



Inputs

## Decoder

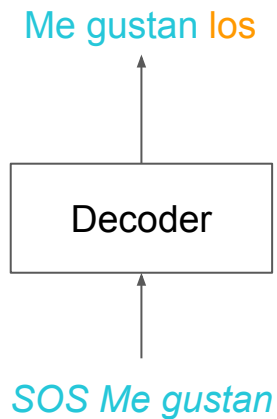


Outputs  
(shifted right)



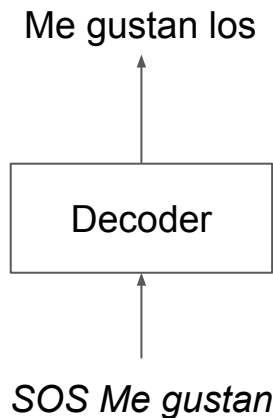
# Training / inference discrepancy

---



# Training / inference discrepancy

---




What decoder knows during inference

*SOS me gustan*

# Training / inference discrepancy

---



What decoder knows during inference

*SOS me gustan*

What decoder knows during training

*SOS me gustan los gatos*

# Masked multi-head attention

---

$$Z_i(Q_i, K_i, V_i) = \text{softmax} \left( \frac{Q_i K_i^T}{\sqrt{d_k}} \right) V_i$$

# Masked multi-head attention

---

$$\frac{Q_i K_i^T}{\sqrt{d_k}} =$$

	SOS	me	gustan	los	gatos
SOS	1.3	0.8	1.3	2.8	2.3
me	2.4	2.8	2.3	6.8	1.9
gustan	1.6	7.4	1.6	0.3	0.5
los	2.1	1.2	9.3	5.2	0.2
gatos	4.3	3.8	6.3	1.8	2.3

# Masked multi-head attention

---

$$\frac{Q_i K_i^T}{\sqrt{d_k}} =$$

	SOS	me	gustan	los	gatos
SOS	1.3	0.8	1.3	2.8	2.3
me	2.4	2.8	2.3	6.8	1.9
gustan	1.6	7.4	1.6	0.3	0.5
los	2.1	1.2	9.3	5.2	0.2
gatos	4.3	3.8	6.3	1.8	2.3

# Masked multi-head attention

---

$$\frac{Q_i K_i^T}{\sqrt{d_k}} =$$

	SOS	me	gustan	los	gatos
SOS	1.3	<del>0.8</del>	<del>1.3</del>	<del>2.8</del>	<del>2.3</del>
me	2.4	2.8	2.3	6.8	1.9
gustan	1.6	7.4	1.6	0.3	0.5
los	2.1	1.2	9.3	5.2	0.2
gatos	4.3	3.8	6.3	1.8	2.3

# Masked multi-head attention

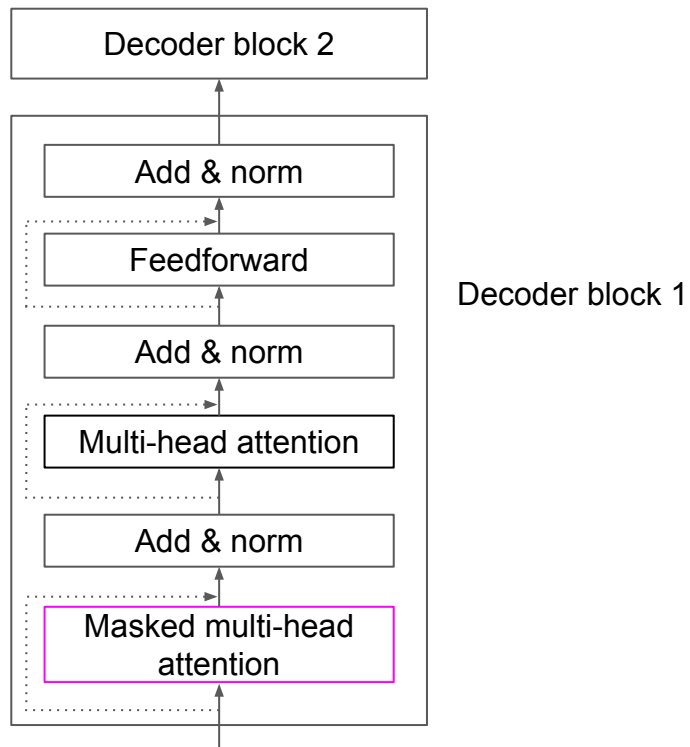
---

$$\frac{Q_i K_i^T}{\sqrt{d_k}} = \begin{array}{c} \text{SOS} \\ \text{me} \\ \text{gustan} \\ \text{los} \\ \text{gatos} \end{array} \begin{array}{ccccc} \text{SOS} & \text{me} & \text{gustan} & \text{los} & \text{gatos} \\ \left[ \begin{array}{ccccc} 1.3 & -\infty & -\infty & -\infty & -\infty \\ 2.4 & 2.8 & -\infty & -\infty & -\infty \\ 1.6 & 7.4 & 1.6 & -\infty & -\infty \\ 2.1 & 1.2 & 9.3 & 5.2 & -\infty \\ 4.3 & 3.8 & 6.3 & 1.8 & 2.3 \end{array} \right] \end{array}$$




# Decoder block

---




# Decoder block

---



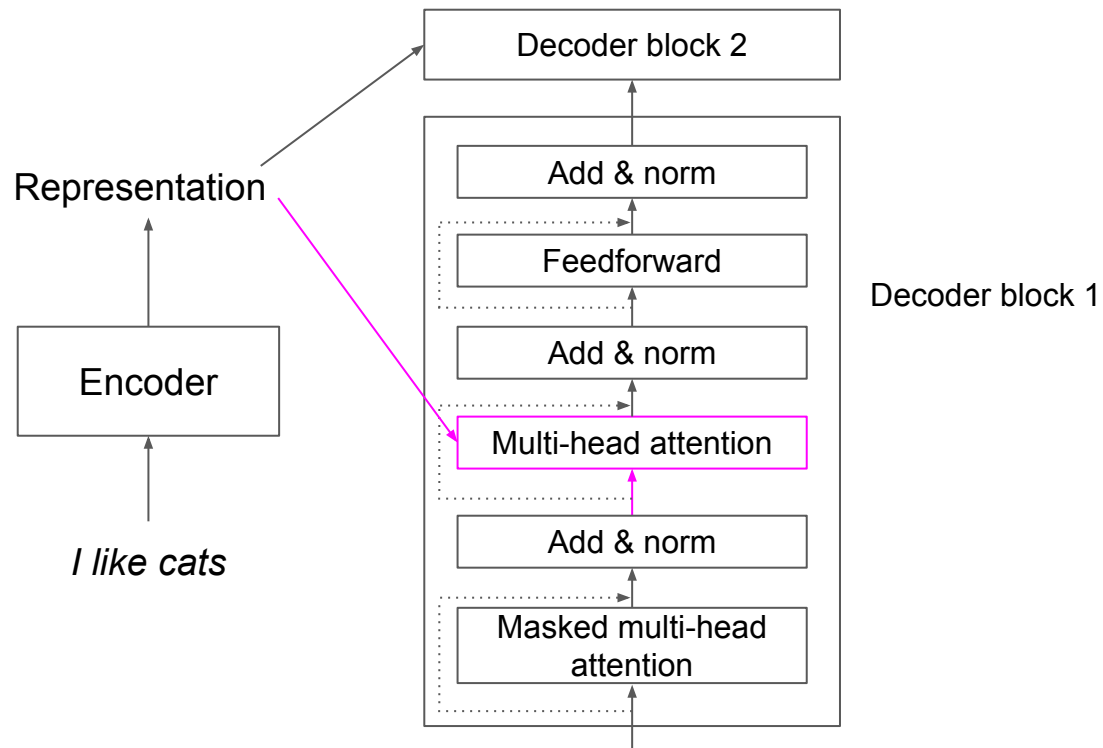
# Decoder block

---




# Decoder block

---



**HOW CAN WE DERIVE  
Q, K, V WITH 2 INPUTS?**



# Deriving Q, K, V

---

- Query matrix (Q) from masked attention input
- Key (K) and value (V) matrices from encoder representation

$$MW_Q = Q$$

$$RW_K = K$$

$$RW_V = V$$

# Deriving Q, K, V

---

- Q holds representation of target sentence
- K, V hold representation of source sentence

**BUT WHY?**



**THIS FEELS SO ARBITRARY**



# Deriving attention matrix

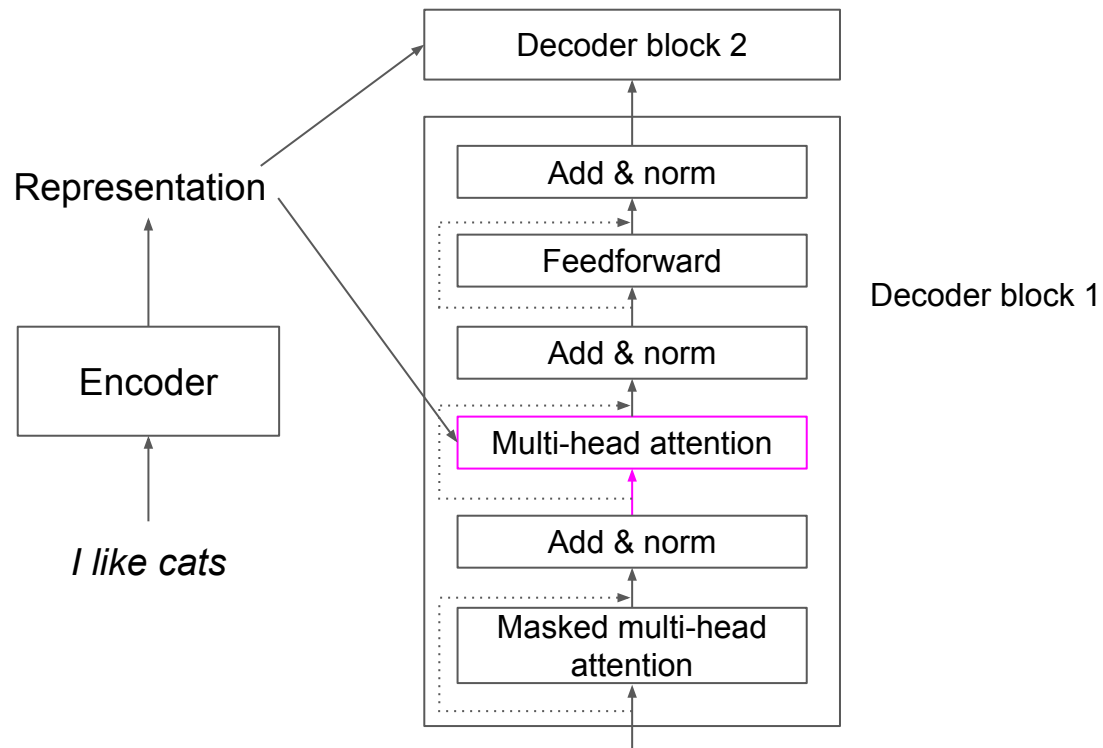
$$Z = \begin{matrix} & \text{I} & \text{like} & \text{cats} \\ \text{SOS} & \begin{bmatrix} 0.7 & 0.2 & 0.1 \end{bmatrix} \\ \text{me} & \begin{bmatrix} 0.6 & 0.3 & 0.1 \end{bmatrix} \\ \text{gustan} & \begin{bmatrix} 0.1 & 0.8 & 0.1 \end{bmatrix} \\ \text{los} & \begin{bmatrix} 0.1 & 0.3 & 0.6 \end{bmatrix} \\ \text{gatos} & \begin{bmatrix} 0.1 & 0.1 & 0.8 \end{bmatrix} \end{matrix} \quad \begin{matrix} \text{I} \\ \text{like} \\ \text{cats} \end{matrix} \begin{bmatrix} 0.4 & 1.0 \\ 1.2 & 2.8 \\ 1.7 & 0.2 \end{bmatrix} \begin{matrix} v_1 \\ v_2 \\ v_3 \end{matrix} = \begin{matrix} \text{SOS} \\ \text{me} \\ \text{gustan} \\ \text{los} \\ \text{gatos} \end{matrix} \begin{bmatrix} \vec{z}_1 \\ \vec{z}_2 \\ \vec{z}_3 \\ \vec{z}_4 \\ \vec{z}_5 \end{bmatrix}$$

$$\vec{z}_3 = 0.1\vec{v}_1 + 0.8\vec{v}_2 + 0.1\vec{v}_3$$

gustan                      I                      like                      cats

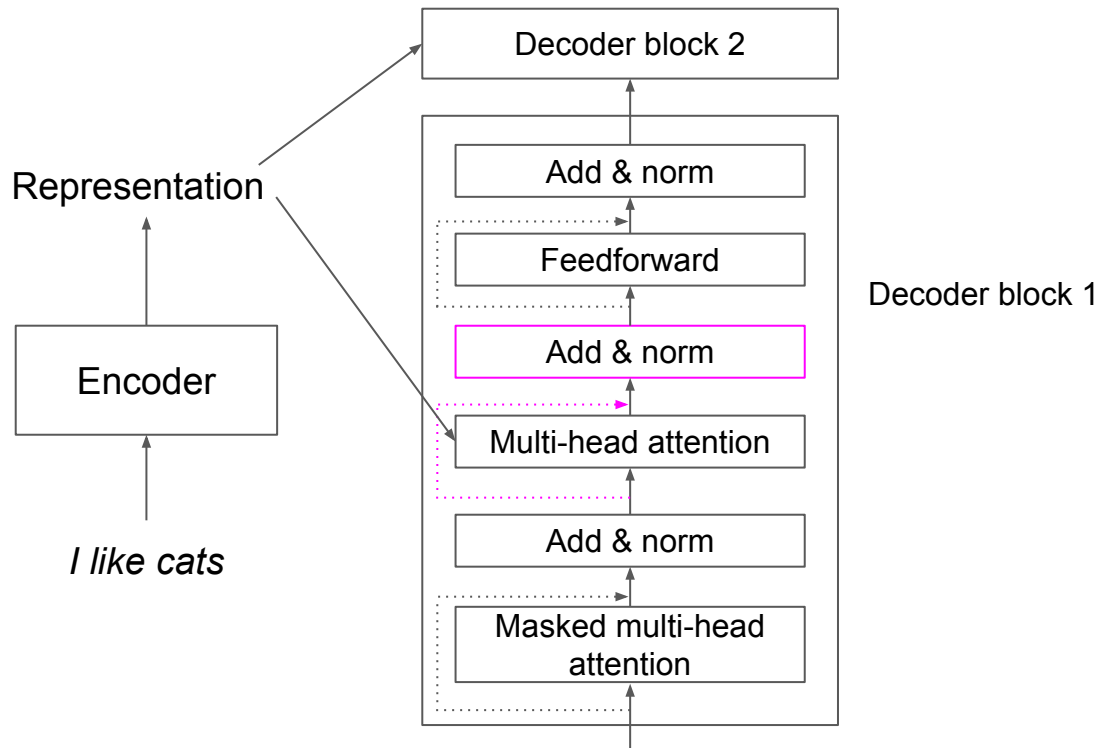
# Decoder block

---



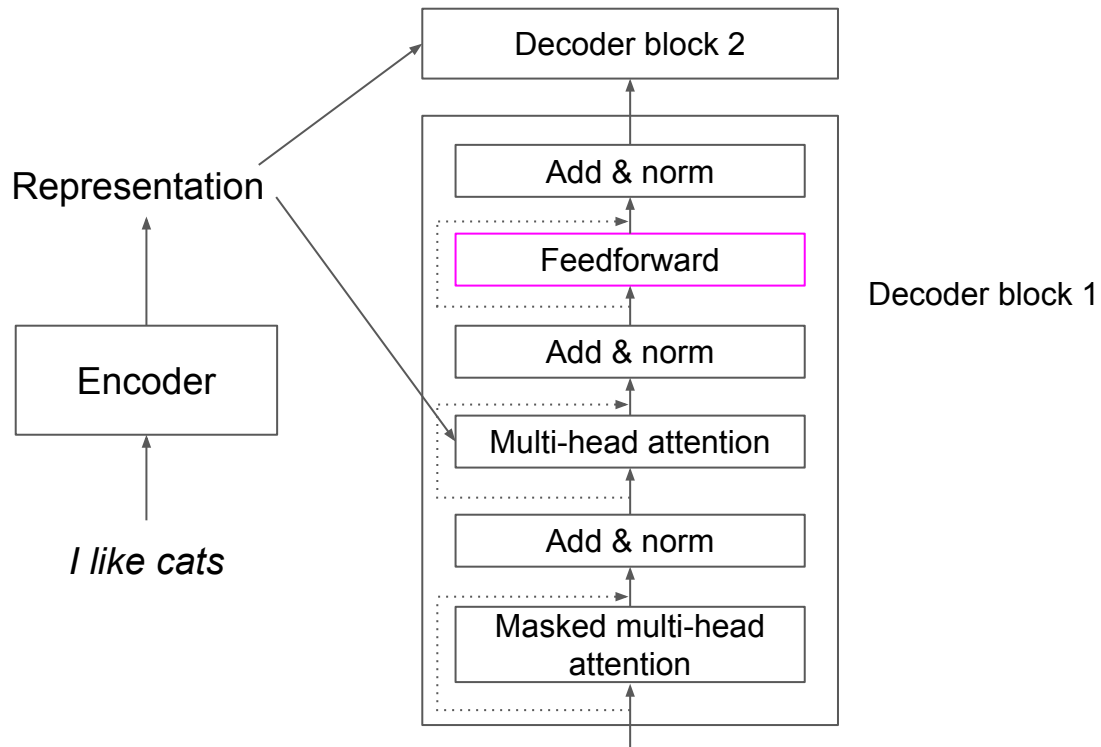
# Decoder block

---



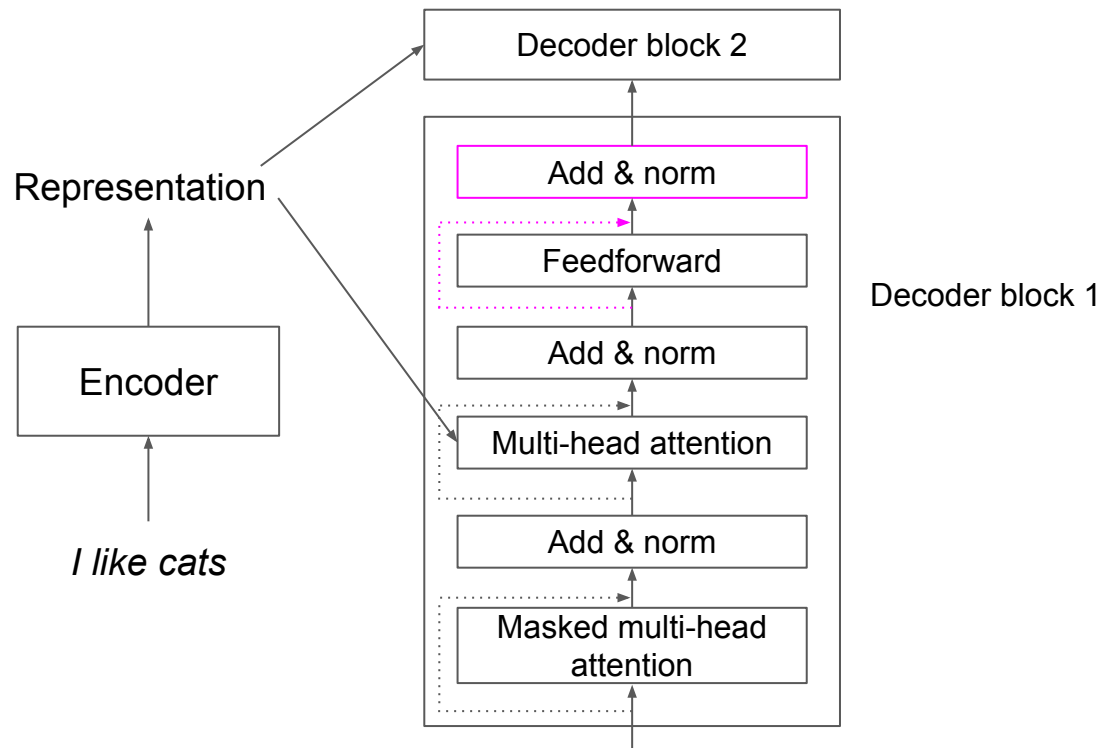
# Decoder block

---



# Decoder block


---



# What's the deeper meaning?


---

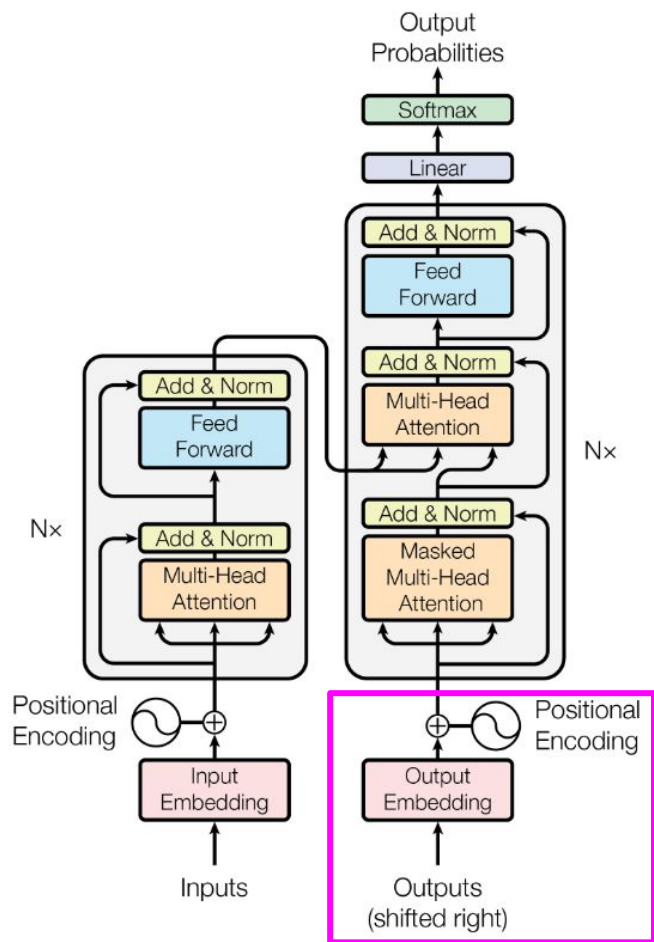
- Masked multi-head attention
- Multi-head attention
- Feedforward
- Add & Norm



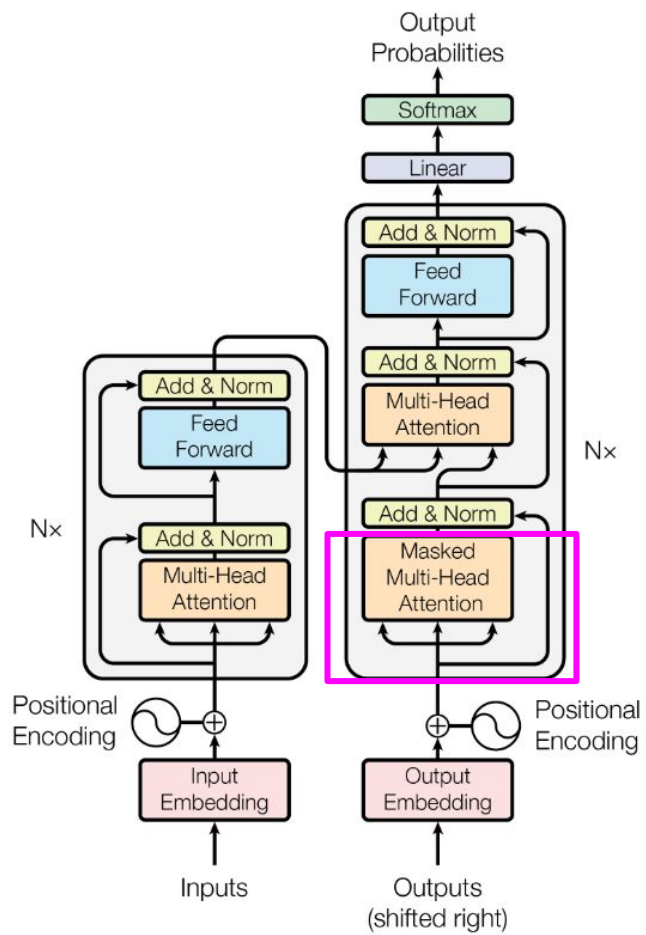
# Linear & softmax layers

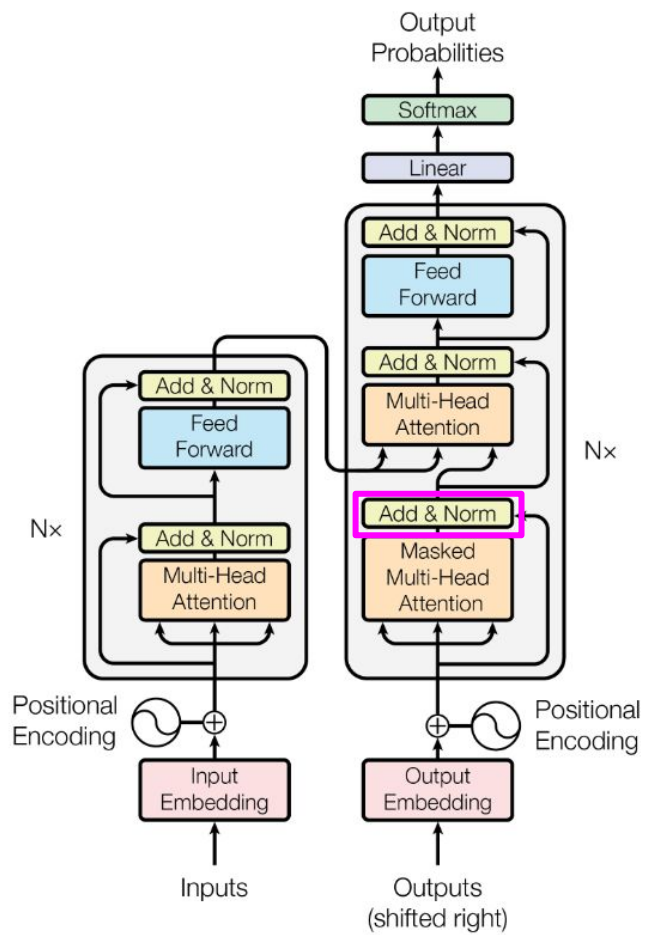
---




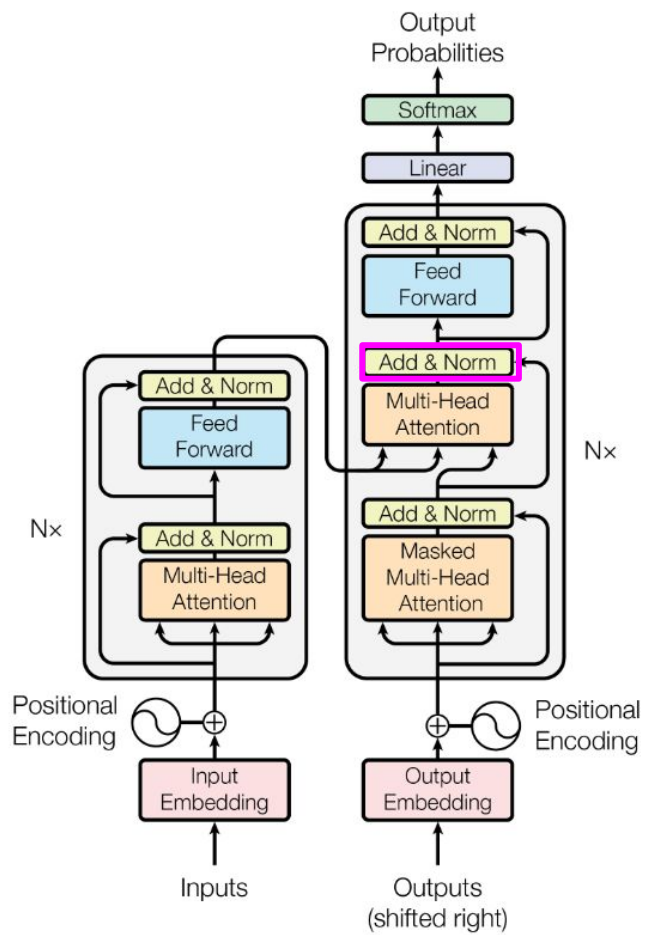


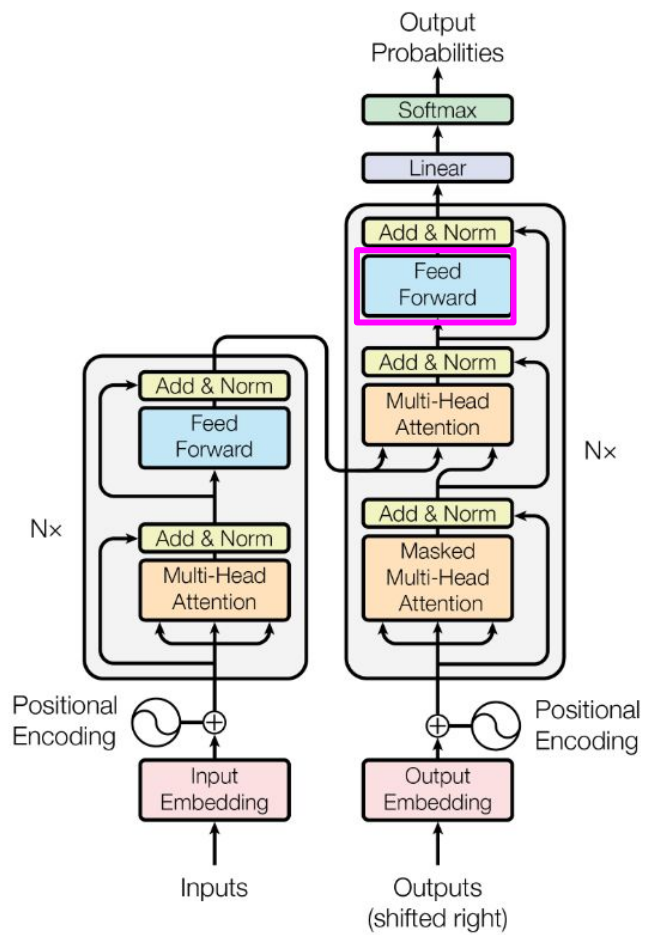


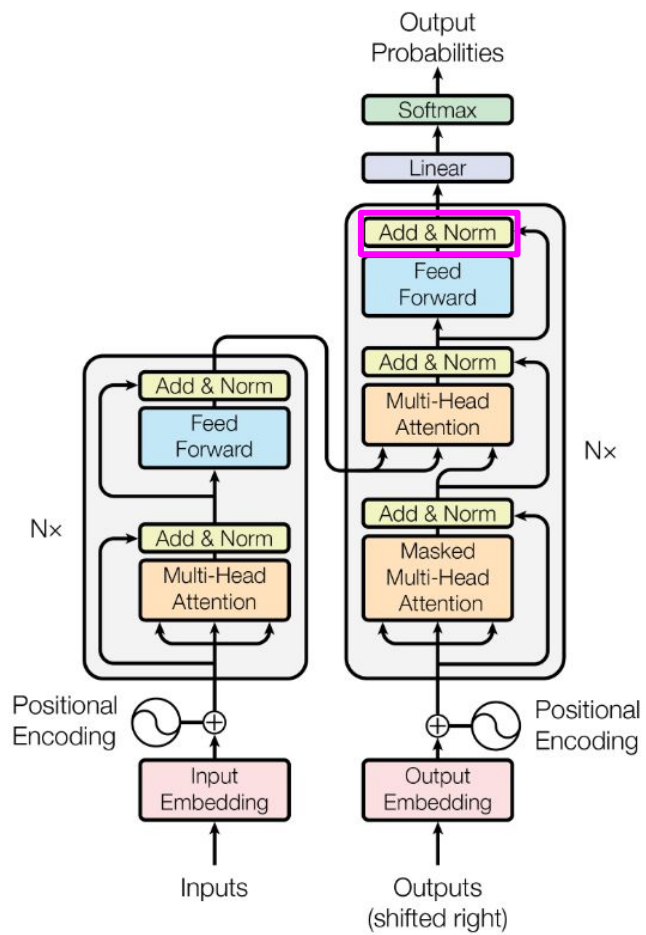













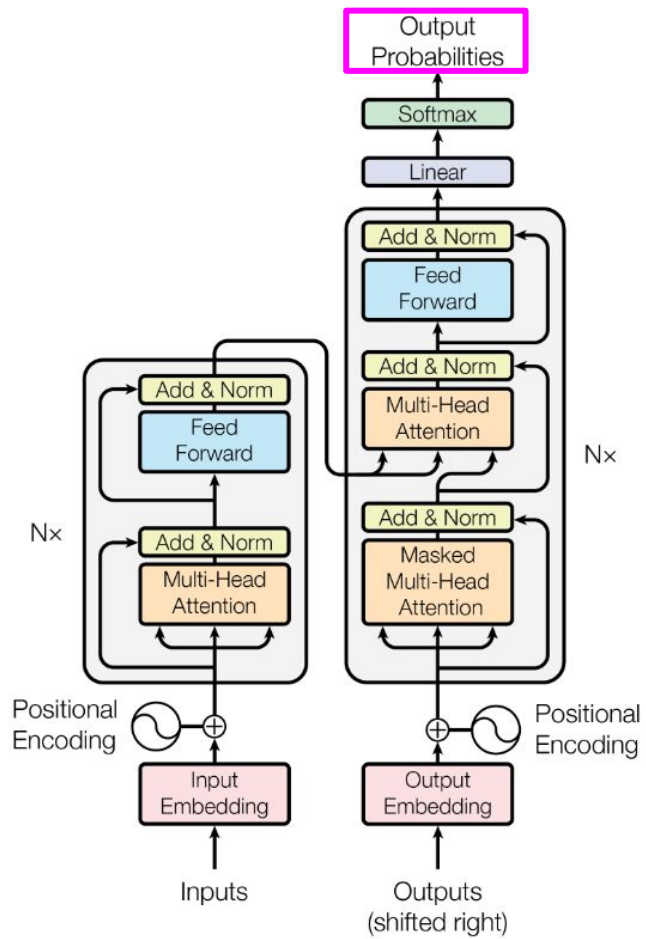












How can you use a  
transformer to  
generate chords?

How can you condition  
generation on  
emotion?

How would you  
evaluate the output of  
a transformer that  
generates melodies?

Do we care about  
overfitting? Do we care  
about perfect  
prediction?

Can transformers  
create truly original  
music? Can they get  
us to transformational  
creativity?

**YOU'VE MADE IT**



**CONGRATS!**



# My experience with Transformers

---

- Most capable model

# My experience with Transformers

---

- Most capable model
- Massive amount of music data needed

# My experience with Transformers

---

- Most capable model
- Massive amount of music data needed
- Music theory bozo


# My experience with Transformers

---

- Most capable model
- Massive amount of music data needed
- Music theory bozo
- Music representation is everything


# Lack of creative control

---



# Lack of creative control

---



**AI ENGINEERS WHO THROW AI AND  
BRUTE-FORCE AT MUSIC GENERATION TASKS**



**IGNORING ALL  
MUSIC KNOWLEDGE**

# My idea

---

- 2-level transformer
- Level 1: Generate high-level music representation
- Level 2: Fill the notes for level 1



# My idea: Music representation

---



# My idea: Level 1 music representation

---

- High-level description of symbolic music

# My idea: Level 1 music representation

---

- High-level description of symbolic music
- Easier to learn
- More coherent generation
- Controllable

# Tips for using Transformers

---

- Get as much (consistent) data as possible

# Tips for using Transformers

---

- Get as much (consistent) data as possible
- Music-informed tokenization

# Tips for using Transformers

---

- Get as much (consistent) data as possible
- Music-informed tokenization
- Transpose data to C / Amin

# Tips for using Transformers

---

- Get as much (consistent) data as possible
- Music-informed tokenization
- Transpose data to C / Amin
- Augment music data

# Tips for using Transformers

---

- Get as much (consistent) data as possible
- Music-informed tokenization
- Transpose data to C / Amin
- Augment music data
- When using pre-trained models:
  - Fine-tune
  - Distillation



**ANY QUESTIONS / DOUBTS/ IDEAS?**



# Activity 1: Bridging symbolic and audio

---

Come up with a strategy / new architecture to repurpose the Transformer architecture for audio-based music gen. What challenges would you face?

Instructions:

- Work in groups (5 people)
- 7' to come up with a solution
- 5' to discuss together

# Museformer

<https://ai-music.github.io/museformer/>

# Activity 2: Transformers go long

---

Long-term coherence has long been a problem in gen mus. Skim through the *Museformer* paper and answer the questions:

- What's the model about?
- How does it try to address long-term coherence?

Instructions:

- Work in pairs
- 10' to read / study
- 5' to discuss together

# Activity 3: LLMs generate music

---

Use ChatGPT (or any other LLM) to generate music. Try to sonify it.

- Is it any good? What surprised you (good / bad)?
- What are its shortcomings?
- How could you improve the generation?

Instructions:

- Work in groups (5 people)
- 10' to play around
- 5' to discuss together

# Assignment 4: Transformer training

---

Train Transformer on Irish folk tunes dataset.

Deadline: 25 January at midnight